



An Alternative Paradigm for Machine Learning

A Detailed Description of the Ryskamp Machine Learning System

By Jeff Horton & Rix Ryskamp

Last Updated October 12, 2017

use**A**ible



Executive Summary

The field of machine learning contains many powerful models that have accomplished some impressive feats. The majority of these machine learning systems are based upon a mathematical model and computations of equations, statistics, and other mathematical devices. While impressive, these traditional math-based systems have some significant drawbacks.

useAible's Ryskamp Learning Machine (RLM) creates a new paradigm for machine learning where a combination of mathematical and logical functions co-exist to create a framework that broadens the capabilities of machine learning in general. This framework is faster, capable of solving more types of problems, completely transparent about the decisions it makes, and adds many other new dimensions to machine learning.

This document will compare the traditional approach to machine learning with the RLM to help readers understand how the different approach to machine learning operates. Although many alternatives exist, for our comparison, we will focus on artificial neural networks, as they are the most widely used at this time. This summary will look at the largest drawbacks with traditional machine learning and compare how the RLM overcomes these drawbacks.

Heavy Computational Requirements

Neural networks contain many neurons that are patterned after the neural network in the human brain. Each of these neurons contains connections to inputs, outputs, or other neurons. Associated with each incoming connection on the neuron is a weight that increases or decreases the significance of a connection to the neuron. For example, a weight of 1 would pass values directly through, while a weight of .5 would cut values in half as they pass through the neuron. A variety of formulas can be used to compute these values and their weights for each neuron, determine its "activation" or lack thereof, and change weight values to perform learning. The nature of traditional neural networks is such that each neuron will converge on its own portion of the problem during the learning cycle. Since neurons do not know in advance which portion of the problem is assigned to them, each neuron must be processed during every learning cycle as any neuron in the network has an equal chance of being relevant to any learning cycle. Thus, each calculation described above will be processed for each neuron in the network once per learning cycle.

This fundamental method creates an exponential computational requirement. The number of neurons required, particularly in what is known as hidden layers, expands as the complexity of the problem expands. As problem complexity increases, the width (number of layers) and depth (number of neurons per layer) increases exponentially.

For example, let's look at The Shape Network, which is solving a problem where the network is passed shapes and colors and must take a different action for each shape/color combination. If the network is known to only need to process three shapes and three colors, a *minimum* of nine



(3x3) neurons are required. When the number of colors and shapes moves to 1,000 possible values for both shape and color, our *minimum* requirement moves to 1,000,000 (1,000x1,000). If the designer does not know the limitation of shapes and colors she will need to add additional neurons to handle any inputs that may occur, resulting in a large number of neurons that may or may not be used. Conversely, if an overall pattern exists between color/shape combinations that causes overlapping solutions among neurons then she may choose to use fewer neurons. This example is not meant to be complete, only to demonstrate that the number of neurons in a network grows exponentially as the problem size and complexity grows. Again, each of these neurons and their associated activation and learning functions must be calculated individually with each pass or cycle of learning.

This computational requirement is demonstrated by the hardware industry's effort to keep up with today's machine learning problems. For example, the NVIDIA DGX-1 V100 has 8 Tesla V100 GPUs, 40 Intel Xeon cores, 40,960 CUDA cores, and 5,120 tensor cores. The need for these various and numerous processors is a direct result of the above explained neuron calculation requirements.

With the Ryskamp Machine Learning engine (RLM), the paradigm changes. Let's look at The Shape Network example with 1,000 colors and 1,000 shapes. The RLM would still require 1,000,000 neurons. The difference is in *how* the neurons are processed. In the RLM neurons are preassigned to a specific portion of the problem (technically a unique input combination). So, the RLM would see a red square, find the neuron associated with red and square and process *ONLY* that neuron in the learning cycle. This means roughly 999,999 fewer calculations are required per learning cycle.

Imagine how much information could be processed with that same NVIDIA DGX-1 with a 1,000,000 times lighter load. Now consider that a set of 1,000,000 neurons is nowhere near the maximum requirement for today's machine learning problems. Training that now takes months or years can be reduced to minutes and days.

Limited to Math Based Problem Solving

The vast majority of machine learning is using an overall mathematical equation or statistical analysis to perform machine learning. Not surprisingly, the problems that these systems are good at solving can be generally explained by mathematics. Machine vision, speech analysis, and providing search engine results are all examples of real world problems that can be explained by math. Each of these examples use aggregation of information to find patterns and therefore solutions. A good way to simplify this concept is by using the phrase "problem x can be solved by statistical analysis of patterns in large sets of data". Search engines can find results by statistical analysis of patterns in large sets of data. Speech analysis is performed by statistical analysis of patterns in large sets of data. Problems where machine learning excels can typically fit well into this phrase.



So, do most problems that face today's organizations fit in this phrase? This is where the use of mathematics without logic breaks down. Many problems are better expressed with logical statements. Computer science often uses "IF, THEN" statements. In these statements an expression is solved and a computer will take one of two paths based upon the result. Most of the software we use today is fundamentally based upon the computation of expressions. However, it is critical to note that in a computer these expressions can be mathematical, logical, or both. For example, "IF $X+Y^2 > Z$ AND A DOES NOT EQUAL 2 THEN . . .".

If we were to remove logical operators from computers, programs and their developers would spend an inordinate amount of time trying to use mathematical expressions to simulate logic that could have been easily expressed by a logical statement. It is the combination of logical operations and mathematical operations that make computers powerful.

At the electronic level computers use combinations of transistors to perform simple AND, OR, XOR, and NOT operations. These operations can be combined to solve logical problems, such as "True AND True = True" or "True OR False = True". These operations can also be combined into algorithms that simulate mathematical functions such as multiplication, division, addition, subtraction, etc. So, we learn from computers that the combination of logic and math are where true problem solving power is born.

The RLM uses this underlying concept within the context of machine learning. The RLM uses neurons, called Smart Neurons[®], that use logic first and then combine mathematics with that logic. Let's consider The Shape Network example. Now let's add switches and dials to our example. Switches will represent logical operations since they can only be on or off and dials will represent mathematical operations as they have an infinite number of positions. The weights in a neural network are very analogous to a dial while the RLM's mechanism is very analogous to a combination of switches and dials. So, if the RLM is sent a red square it will assign a Smart Neuron[®] to [red], [square]. Let's say that our example must produce an answer of yes or no. The RLM will see red square for the first time and produce a result (we will discuss how later). If our result is yes for a red square and our network receives negative feedback (receiving positive or negative feedback is always part of machine learning) it will immediately store a switch for [red], [square], [yes] is [negative]. Next time we receive a red square, the RLM produces a no and receives positive feedback. It will now store a switch for [red], [square], [no] is [positive]. These two cycles are all that are required to correctly make predictions for red square. Unless something else is added to this problem, the correct switch for red square will always be used to produce a no.

Since our traditional neural network is using only dials, the neurons will have to see this problem many times before the dials "converge on the problem", that is they find a combination of dials that produces correct results for all cases. This is inherent, since dials have an almost infinite number of positions, while switches are simply on or off.



Now let's look at an actual proof of this underlying concept. At <http://useaible.com/machine-learning-challenge/> you will find a simple maze puzzle that allows the user to run maze with the RLM against many popular machine learning systems including implementations in Google's TensorFlow®. A quick review of this challenge shows that after one pass through the maze, the RLM will get a perfect score every subsequent time. No other system or method has been able to solve the maze (even after running for months). While it is certainly likely that one could set up a combination of traditional methods to learn to solve this maze, the point that it is a huge workaround for math to be solving a logical problem remains. The RLM can solve the maze simply as a result of saving the switches mentioned above. For example, [X=5], [Y=2], [Direction to Move=UP] will always be significantly simpler than a series of dials that tries to randomly converge upon this same result. Now consider that this random convergence of dials would be multiplied by the number of locations on the maze.

While this makes sense for logical only problems, we have not yet addressed mathematical and logical problems. We will go back to The Shape Network example. Logic is an easy solution when red is either present or not. What about shades of red? How will the RLM predict correctly if a new shade of red is presented? Let's change our input that was called "color" to "shade of red" and say that it can be a value between 1 and 1000. We will also assume that there is a clear linear pattern to this number, meaning that shades 99 and 100 look almost identical while 10 and 500 would not appear to be close shades of red.

We now present our RLM with square 500. The RLM produces a "yes", receives positive feedback, and saves a switch for [square], [500], [yes], [positive]. Then we present square 600. The RLM produces a "no", receives positive feedback, and saves a switch for [square], [600], [no], [positive].

When the RLM subsequently receives square 525 what will it output? It would be poor practice to save a switch for every one of the 1,000 shades of red before being able to make a prediction. This is where the RLM's equivalent of dials comes into play. A dial, called a linear bracket in the RLM, may be set by default to 10% for shade of red. Since the 10% represents a tolerance value, any value between 450 and 550 will point to [square], [500], [yes], [positive] until additional Smart Neurons® appear in that range. As more Smart Neurons® appear, learning becomes more specific and less generalized. For example, when the 525 above is used it will produce a yes value. If the feedback system tells us that yes was negative then we now know that even though 500 and 525 are close, they are not to be treated the same. This process, called maturing of neurons, allows for an unprecedented combination of categorization and specific learning wherein an "edge case" that is an unexpected value within a certain category is tracked separately from the category around it and handled perfectly.

This handling of edge cases is critical in many fields. For example, in medicine a slight difference in the symptoms on stomach pain can mean the difference between giving a patient a pain reliever or sending her to immediate surgery for appendicitis. Traditional neurons



networks are not well suited to find edge cases within broader categories without extensive repetition and training on large neural networks.

Complicated

Traditional machine learning systems are complicated compared to other computer science disciplines. Many machine learning engineering positions require a computer science degree and extensive post graduate work specifically in a machine learning field. This high threshold of required knowledge combined with an expansion of the use of machine learning by organizations has created a shortage of talent and driven up the costs associated with machine learning.

Although a comprehensive look at the reasons for these trends are outside the scope of this paper, we will look at a few factors that contribute to this trend. Firstly, there are many types of machine learning in addition to subtypes of formulas, parameters, and other configuration steps. Matching a problem to the correct type of machine learning with the correct settings is a significant task that requires in-depth training and skills. Secondly, the large dependence on relatively complex mathematics requires users to have a strong background in mathematics.

When compared to most other computer science fields, machine learning is behind the curve for simplifying user experiences and abstracting complicated concepts into simplified encapsulations. For example, a software developer no longer needs to understand the complexities associated with storage and retrieval of information to open a file in a computer program. She does not need to know what type of storage is being used, whether or not the file is local or accessed via a network, or any other myriad complexities associated with file access. These complexities have been abstracted into a simple object that can be accessed through simple commands with no understanding about or training on the underlying systems.

This abstraction has not been popularized with machine learning. There are many reasons that contribute to this situation; however, as long as the user of a system must be educated enough to set every variable perfectly for something to work correctly such encapsulation is not possible. In traditional machine learning, something as simple as the wrong loss function or the number of hidden layers can have drastic negative impacts that may be costly and time consuming to correct.

The RLM changes this model. Its “mostly logic plus some math” approach takes the bulk of complexity out of the system. There are simply less moving parts. The moving parts that exist are primarily abstracted into simple APIs for use by common software developers. If a user wishes to be an RLM expert, she may do so with much less effort than would be required to master a traditional counterpart machine learning system. The underlying concepts that drive the RLM may be taught in weeks instead of months and mastered in months instead of years. UseAble frequently teaches people with no previous machine learning experience to use the RLM to solve real world problems.



Not Explainable

Traditional neural networks cannot explain why they make the decisions they make. This problem, often called the “black box” problem, is a serious hindrance to the adoption of machine learning in critical situations where explainability is paramount. Examples include financial, medical, transportation, industrial, and other industries where logic must be scrutinized before it can be used.

A neural network’s lack of explainability is easy to understand as is the RLM’s ease of explainability. Let’s go back to our example from above and think about the dials. Let’s present our standard neural network with a red square. Until the feedback systems tells the network that the response for red square is perfect, it will continue to move all or some of the dials in the neurons present in the system until it achieves its desired result. No record is kept of the dial positions over time. Moreover, even if you did track knob position over time, the statistical nature of the reason why a dial was turned mathematically sound yet not understandable by humans trying to find a cause and effect relationship. We call this compound aggregation.

When you say that the average of 4, 5, and 6 is 4, you have aggregated 4, 5, and 6 into 4. Now let’s say that the result of another aggregation tells us to make our 4 into a 5 as it deems 5 to now be more correct. Now we will take our 5 and combine with other results, say 10, 15 and average those to 10. Then we are told by another aggregation to move the 10 to 9 as it now deems 9 to be more correct. Now imagine this process taking place billions of times during a training session. Now imagine that your boss walks in and asks you why the final answer is 18 when his experience says it should be 2. This process is analogous to the everyday experience of machine learning engineers. When the answers are what the boss expects it is magic. When they are not the boss wants to start a discussion about manpower and hardware costs.

No complex mathematical formulas means no compound aggregation. The “switches” we saved above (e.g. [red], [square], [yes], [positive]) are all stored in a database along with time and date information and relationships between all inputs ever captured and all decisions ever made. With our visualization tools you can tell your boss exactly why the network chose 18. If he disagrees you can walk through the feedback logic to see if the machine was told the wrong answer or if it figured out something the boss did not already know. This entire process takes minutes and is clearly understandable. On our GitHub page at <https://github.com/useaible/ryskamplearningmachine> there is a sample application called RetailPoCSimple that demonstrates a very simple version of the explainability visualizations.

Conclusion

The RLM presents us with an alternate paradigm for solving problems with machine learning. While the RLM could be applied in many circumstances, we recommend considering it for problems where traditional machine learning has not already excelled, problems where resources and return on investment are critical, and problems where traditional machine learning is not an option and you desire the benefits of a self-learning system over traditional



algorithmic programing, including costs, time-to-market, ROI, and optimization beyond human abilities.